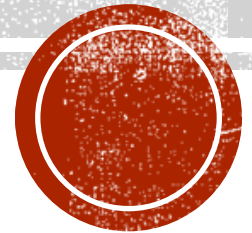


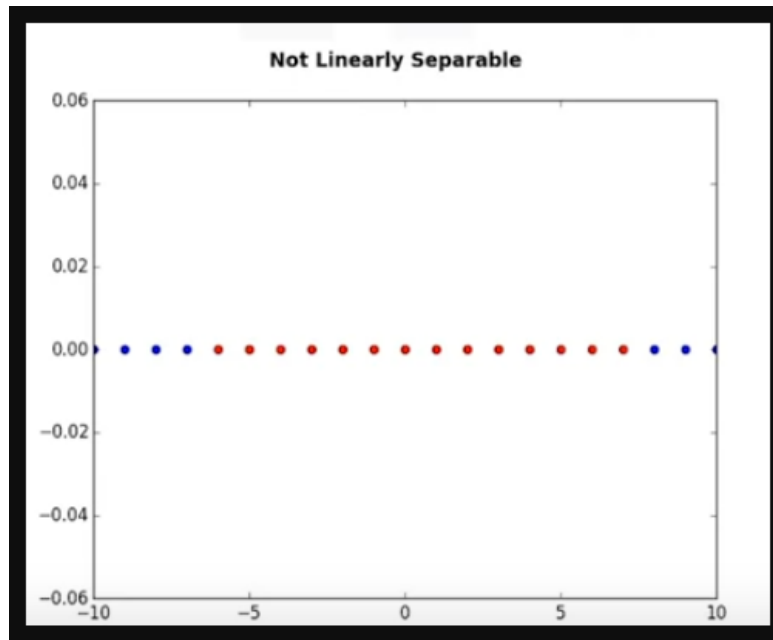
# SUPPORT VECTOR MACHINES

Dr. Brian Mc Ginley



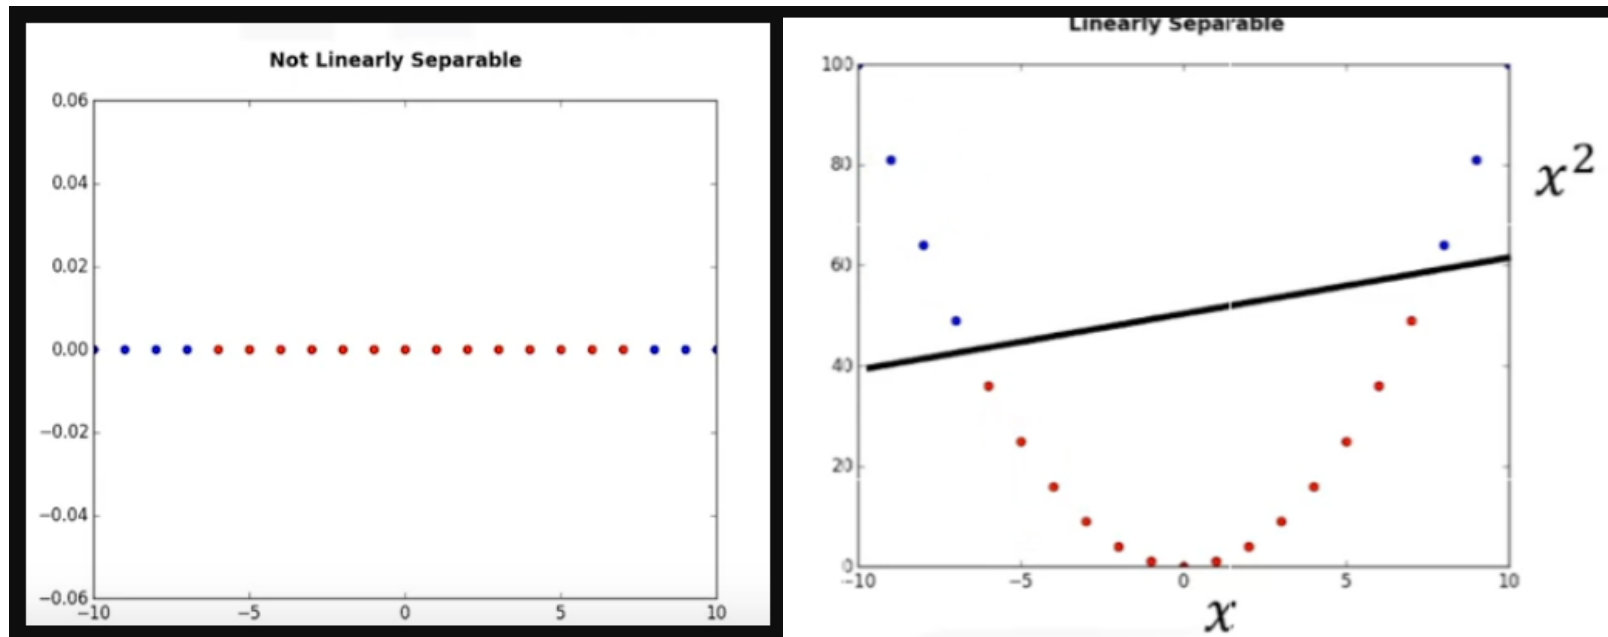
# SVM - TWO KEY IDEAS

- Assuming linearly separable classes, learn separating hyperplane with maximum margin (SVC)
- But what if data is not linearly separable



# SVM - TWO KEY IDEAS

- Assuming linearly separable classes, learn separating hyperplane with maximum margin (SVC)
- Expand input into high-dimensional space to deal with linearly non-separable cases



# SUPPORT VECTOR CLASSIFIER: HARD MARGIN

- Finally, formulate our optimization problem: Find a decision boundary that maximises the distance to both classes – i.e. maximises the margin,  $M$ , while maintaining zero misclassifications

$$\begin{cases} \max_w \frac{2}{\|w\|} \\ \text{such that } y^{(i)}(w^T x^{(i)} + w_0) \geq 1, & \forall i \end{cases}$$

- Maximising  $\frac{2}{\|w\|}$  is the same as minimizing  $\|w\|$ . However L2 optimisations are more stable. Therefore:

$$\begin{cases} \min_w \|w\|^2 \\ \text{such that } y^{(i)}(w^T x^{(i)} + w_0) \geq 1, & \forall i \end{cases}$$

- This is a quadratic optimisation problem, has linear constraints and there is a unique solution.
- Calculus again! (Lagrange multipliers if you want to look up the maths)



# SOFT MARGIN SOLUTION

- To relax the constraints, our problem is re framed as


$$\begin{cases} \min_w \|w\|^2 + C \sum_{i=1}^N \xi_i \\ \text{such that } y^{(i)}(w^T x^{(i)} + w_0) \geq 1 - \xi_i \text{ and } \xi_i \geq 0, \end{cases} \quad \forall i$$

- C is a **regularisation** parameter: (some notes will use  $\lambda$  instead of C , sklearn uses C)
  - Small C allows constraints to be easily ignored  $\rightarrow$  large margin
  - Large C makes constraints hard to ignore  $\rightarrow$  narrow margin
  - $C \rightarrow \infty$  enforces all constraints: hard margin
- This is still a quadratic optimization problem and there is a unique minimum. Note: there is only one parameter, C (that you choose/cross-validation).
- In general, the best C parameter depends on the situation. Experiment (Cross-Validation). One note: larger C takes more computation to train.




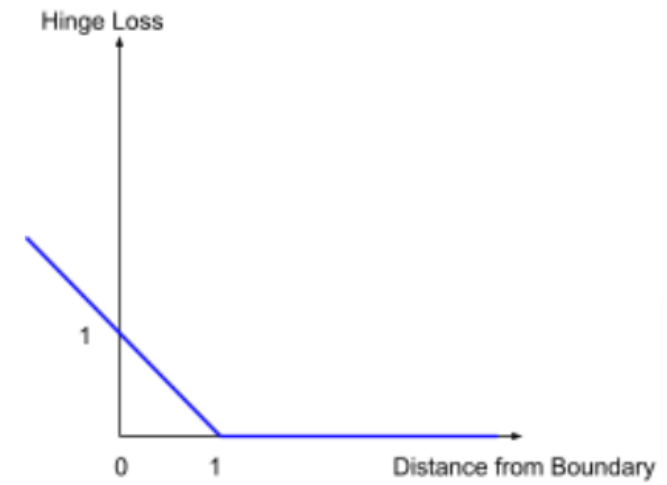
# SLACK VARIABLES

- We can add a variable  $\xi_i \geq 0$  (must be bigger than 0) for each point/sample.
  - For  $0 < \xi \leq 1$  point is between margin and correct side of hyperplane. This is called a **margin violation**
  - For  $\xi \geq 1$  point is **misclassified**
  - For  $\xi = 0$  point is the correct side of the margin.



# SLACK VARIABLES

- We can add a variable  $\xi_i \geq 0$  (must be bigger than 0) for each point/sample.
  - For  $0 < \xi \leq 1$  point is between margin and correct side of hyperplane. This is called a **margin violation**
  - For  $\xi \geq 1$  point is **misclassified**
  - For  $\xi = 0$  point is the correct side of the margin.



# HINGE LOSS

- The hinge loss function for a single data point is defined as:

$$L(y^{(i)}, f(x^{(i)})) = \max(0, 1 - y^{(i)}f(x^{(i)}))$$

- And in the case of SVM as:

$$L(y^{(i)}, f(x^{(i)})) = \max(0, 1 - y^{(i)}(w^T x^{(i)} + b))$$





# HINGE LOSS

- The hinge loss function for a single data point is defined as:

$$L(y^{(i)}, f(x^{(i)})) = \max(0, 1 - y^{(i)}f(x^{(i)}))$$

- And in the case of SVM as:

$$L(y^{(i)}, f(x^{(i)})) = \max(0, 1 - y^{(i)}(w^T x^{(i)} + b))$$

$y^{(i)}$	$(w^T x^{(i)} + b)$	$1 - y^{(i)}(w^T x^{(i)} + b)$	$\max(0, 1 - y^{(i)}(w^T x^{(i)} + b))$
1	1	0	0
1	2	-1	0
-1	-0.5	0.5	0.5
-1	4	5	5



# HINGE LOSS

- The hinge loss function for a single data point is defined as:

$$L(y^{(i)}, f(x^{(i)})) = \max(0, 1 - y^{(i)} f(x^{(i)}))$$

- And in the case of SVM as:

$$L(y^{(i)}, f(x^{(i)})) = \max(0, 1 - y^{(i)}(w^T x^{(i)} + b))$$

- Rewriting the above:

$$\xi_i \geq 1 - y^{(i)}(w^T x^{(i)} + b) \text{ and } \xi_i \geq 0$$

$$y^{(i)}(w^T x^{(i)} + b) \geq 1 - \xi_i \text{ and } \xi_i \geq 0$$



# PRIMAL FORM

- Here's the classifier:

$$f(x) = w^T x + b$$

- Here's the optimisation problem:

$$\begin{cases} \min_w \|w\|^2 + C \sum_{i=1}^N \xi_i \\ \text{such that } y^{(i)}(w^T x^{(i)} + w_0) \geq 1 - \xi_i \text{ and } \xi_i \geq 0, \end{cases} \quad \forall i$$

- However, it can be written another way (it has loads of ways of being written). This next way should be the "most" useful way.



# DUAL FORM

- If you formulate the linear classifier instead as (instead of <sup>(i)</sup>, I am going to write it as a subscript  $\alpha_i$  for a particular sample):

$$f(x) = \sum_{i=1}^N \alpha_i y_i (x_i^T x) + b$$

*Remember (Primal)*  
 $f(x) = w^T x + b$

- we now solve an optimisation problem over  $\alpha_i$  . i.e.

$$\max_w W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j=1}^N \sum_{k=1}^N \alpha_j \alpha_k y_j y_k (x_j^T x_k)$$

Subject to  $0 \leq \alpha_i \leq C$  for all  $i$  and  $\sum_i \alpha_i y_i = 0$



# DUAL FORM

- To classify a new point/sample:

- **Primal** version of classifier:

$$f(x) = w^T x + b$$

- Dual version of classifier

$$f(x) = \sum_{i=1}^N \alpha_i y_i (x_i^T x) + b$$

- and take the sign as before as the classification output.



# DUAL FORM

- To classify a new point/sample:

- **Primal** version of classifier:

$$f(x) = w^T x + b$$

- **Dual** version of classifier

$$f(x) = \sum_{i=1}^N \alpha_i y_i (x_i^T x) + b$$

- and take the sign as before as the classification output.
  - The dual form appears to have the disadvantage of a kNN classifier — it requires the training data points  $x_i$ . However, many of the  $\alpha_i$ 's are zero. The ones that are non-zero define the support vectors  $x_i$ . You could write it as


$$f(x) = \sum_{x_i \text{ is a support vector}} \alpha_i y_i (x_i^T x) + b$$

- then we classify a point by finding the sign of the above.




# NON LINEARLY SEPARABLE DATA

- How do we deal with this?



# NON LINEARLY SEPARABLE DATA

- It is not possible to find a hyperplane or a linear decision boundary for some classification problems. If we project the data into a higher dimension from the original space, we may get a hyperplane in the projected dimension that helps to classify the data.






# A SOLUTION - MAP TO HIGHER DIMENSION


- Data is linearly separable in 3D
- The problem can still be solved by a linear classifier

$$\Phi(x) = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \rightarrow \begin{pmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2}x_1x_2 \end{pmatrix} \quad R^2 \rightarrow R^3$$



# A SOLUTION - MAP TO HIGHER DIMENSION

- Data is linearly separable in 2D
- The problem can still be solved by a linear classifier




$$\Phi(x) = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \rightarrow \begin{pmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2}x_1x_2 \end{pmatrix} \quad \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

# SVM IN A TRANSFORMED FEATURE SPACE

- Task: Learn linear classifier  $\mathbf{w}$  for  $\mathbb{R}^D$ :



$$f(x) = \mathbf{w}^T \Phi(x) + b$$

- Where  $\Phi(x)$  is a feature map




# EXAMPLE: THE XOR PROBLEM

$x_1$	$x_2$	$x_1 \text{ XOR } x_2$
0	0	0
0	1	1
1	0	1
1	1	0



# SVM NON-LINEAR DECISION BOUNDARIES

- The idea: instead of tweaking the definition of SVC to accommodate non-linear decision boundaries, we map the data into a feature space in which the classes are linearly separable (or nearly separable):
  - Transform  $\mathbf{x} \rightarrow \phi(\mathbf{x})$
  - The linear algorithm depends only on  $\mathbf{x}^T \mathbf{x}_i$ , hence transformed algorithm depends only on  $\phi(\mathbf{x})^T \phi(\mathbf{x}_i)$
  - Use kernel function  $K(\mathbf{x}, \mathbf{x}_i)$  such that  $K(\mathbf{x}, \mathbf{x}_i) = \phi(\mathbf{x})^T \phi(\mathbf{x}_i)$



# COMPLEXITY

- After projecting the data into a higher dimension, we could find the hyperplane which classifies the data.
- Usually, the computational cost will increase, if the dimension of the data increases.
- Kernel helps to find a hyperplane in the higher dimensional space without increasing the computational cost much.





# KERNEL TRICK

$$\Phi(x) = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \rightarrow \begin{pmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2}x_1x_2 \end{pmatrix} \quad R^2 \rightarrow R^3$$

- $\Phi$  is a function that transforms  $x$  from 2D to 3D.
- We can now have a decision boundary in this 3D space:

$$w_0 + w_1x_1^2 + w_2x_2^2 + w_3\sqrt{2}x_1x_2 = 0$$

- Project  $x$  into high dim space and calculate dot product

$$\phi(x_i) \cdot \phi(x_j) = (x_{i1}^2, x_{i2}^2, \sqrt{2}x_{i1}x_{i2})(x_{j1}^2, x_{j2}^2, \sqrt{2}x_{j1}x_{j2})$$

- The Kernel Trick:

$$\phi(x_i) \cdot \phi(x_j) = x_{i1}^2x_{j1}^2 + x_{i2}^2x_{j2}^2 + 2x_{i1}x_{i2}x_{j1}x_{j2}$$

- The kernel is defined as:  $K(x_i, x_j) = (x_i^T \cdot x_j)^2$

$$(x_i^T \cdot x_j)^2 = \{(x_{i1}, x_{i2})^T \cdot (x_{j1}, x_{j2})\}^2 = (x_{i1}x_{j1} + x_{i2}x_{j2})^2$$

$$(x_i^T \cdot x_j)^2 = x_{i1}^2x_{j1}^2 + x_{i2}^2x_{j2}^2 + 2x_{i1}x_{i2}x_{j1}x_{j2}$$





# KERNEL TRICK EXAMPLE

- Let's say we have two features  $x_i = (x_{i1}, x_{i2})$ , i.e.  $\in \mathbb{R}^2$ .
- Assume we have some transformation to convert to a four-dimensional feature space  $(x_{i1}^2, x_{i1}x_{i2}, x_{i2}x_{i1}, x_{i2}^2)$ . It requires  $O(n^2)$  time to calculate  $n$  data points in the 4-dimensional space.
- If we want to calculate the dot product in the 4-dim space the standard way it's:
  1. Convert each point from  $\mathbb{R}^2 \rightarrow \mathbb{R}^4$  by applying the transformation

$$\phi(x_i) = (x_{i1}^2, x_{i1}x_{i2}, x_{i2}x_{i1}, x_{i2}^2) \quad \phi(x_j) = (x_{j1}^2, x_{j1}x_{j2}, x_{j2}x_{j1}, x_{j2}^2)$$

2. Dot product the two vectors  $\phi(x_i) \cdot \phi(x_j)$



# KERNEL TRICK EXAMPLE

- Example, say  $x_i = (1, 2)$  and  $x_j = (3, 5)$
- Transform them through  $\phi$  to get:

$$\phi(x_i) = (1, 2, 2, 4)$$

$$\phi(x_j) = (9, 15, 15, 25)$$

- Get the dot product:

$$\phi(x_i) \cdot \phi(x_j) = (1, 2, 2, 4) \cdot (9, 15, 15, 25) = 169$$

- But, the kernel of our function is actually:

$$k(x_i, x_j) = (x_i^T x_j)^2$$

- So, we can calculate with our example directly as:

$$k(x_i, x_j) = ((1, 2) \cdot (3, 5))^2 = (3 + 10)^2 = 169$$

- The standard method of calculating this requires  $O(n^2)$  but kernel requires just  $O(n)$



# KERNEL TRICK

- Kernel helps to find a hyperplane in the higher dimensional space without increasing the computational cost much.
  - <https://www.quora.com/What-is-the-kernel-trick/answer/Chitta-Ranjan-2> I think is a really good detailed answer about this "Kernel Trick".
- Since the feature space  $R^D$  is extremely high dimensional, computing  $\Phi$  explicitly can be costly. This is because if  $D \gg d$  then there are many more parameters to learn for  $w$ .
  - Instead, we note that computing  $\Phi$  is unnecessary.

- Classifier:

$$f(x) = \sum_{i=1}^N \alpha_i y_i (x_i^T x) + b$$
$$\rightarrow f(x) = \sum_{i=1}^N \alpha_i y_i (\phi(x_i)^T \phi(x)) + b$$



# KERNEL TRICK

- Optimisation:

$$\max_w W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j=1}^N \sum_{k=1}^N \alpha_j \alpha_k y_j y_k (x_j^T x_k)$$

$$\rightarrow \max_w W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j=1}^N \sum_{k=1}^N \alpha_j \alpha_k y_j y_k \left( \phi(x_j)^T \phi(x_k) \right)$$

Subject to  $0 \leq \alpha_i \leq C$  for all  $i$  and  $\sum_i \alpha_i y_i = 0$

- Note that  $\Phi(x)$  only appears in pairs  $\Phi(x_j)^T \Phi(x_i)$ . So we only need to compute these.
- Once these products are computed (dot product, scalar product), only the  $N$  dimensional vector  $\alpha$  needs to be learnt.
- Write  $\Phi(x_j)^T \Phi(x_i) = k(x_j, x_i)$ . This is known as the **Kernel**



# KERNEL TRICK

- Classifier:

$$f(x) = \sum_{i=1}^N \alpha_i y_i k(x_i, x) + b$$

- Optimisation:

$$\max_w W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{j=1}^N \sum_{k=1}^N \alpha_j \alpha_k y_j y_k k(x_j, x_k)$$

Subject to  $0 \leq \alpha_i \leq C$  for all  $i$  and  $\sum_i \alpha_i y_i = 0$



# KERNEL TRICK SUMMARY

- Now:  $\phi(x) = \begin{pmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2}x_1x_2 \end{pmatrix} \quad R^2 \rightarrow R^3$
- Can calculate  $\phi(x)^T \cdot \phi(z) = \dots = (x^T z)^2$
- Classifier can be learnt and applied without explicitly computing  $\Phi(x)$ .
- All that is required is the kernel  $k(x, z) = (x^T z)^2$
- Complexity of learning depends on  $N$  (the size of the training set) not on  $D$  (the higher dimension). This is a lot faster.



# SVM KERNELS

- Linear kernels (kernel='linear')  $k(\mathbf{x}_1, \mathbf{x}_2) = \mathbf{x}_1^T \mathbf{x}_2$
- Sigmoid Kernel (kernel='sigmoid')  $k(\mathbf{x}_1, \mathbf{x}_2) = \tanh(\kappa \mathbf{x}_1^T \mathbf{x}_2 + \theta)$ 
  - $\kappa$  and  $\theta$  are hyperparameters.



# SVM POLYNOMIAL KERNEL

- Polynomial kernels (kernel='poly')  $k(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1^T \mathbf{x}_2 + c)^d$  for any  $d > 0$ 
  - Contains all polynomial terms up to degree  $d$ .
  - $d$  is a hyperparameter.
  - $c$  is also a hyperparameter that varies the importance of the non-linear terms
- $c$  (called coef0 in sklearn) – default is 0

$c$ Value	Expanded Polynomial Terms
$c = 0$	$(x_1 y_1)^2 + (x_2 y_2)^2 + 2(x_1 y_1)(x_2 y_2)$
$c = 1$	$(x_1 y_1)^2 + (x_2 y_2)^2 + 1 + 2(x_1 y_1)(x_2 y_2) + 2(x_1 y_1) + 2(x_2 y_2)$
$c = 10$	$(x_1 y_1)^2 + (x_2 y_2)^2 + 100 + 2(x_1 y_1)(x_2 y_2) + 20(x_1 y_1) + 20(x_2 y_2)$


- As  $c$  increases, the constant term dominates more, thereby smoothing the decision boundary.
- The linear terms ( $x_1 y_1$  and  $x_2 y_2$ ) grow proportionally to  $c$ , overshadowing higher-order interactions for large  $c$ .





# SVM RBF KERNEL

$$f(x) = a \exp\left(-\frac{(x-b)^2}{2c^2}\right)$$



- Gaussian kernels (kernel='rbf')  $k(\mathbf{x}_1, \mathbf{x}_2) = \exp\left(-\frac{\|\mathbf{x}_1 - \mathbf{x}_2\|^2}{2\sigma^2}\right)$  for  $\sigma > 0$ .
  - Infinite dimensional feature space for Radial Basis Function kernel
  - $\sigma$  is a hyperparameter.
- In sklearn , the RBF kernel is expressed in terms of  $\gamma$  instead of  $\sigma$ :
  - $k(\mathbf{x}_1, \mathbf{x}_2) = \exp(-\gamma\|\mathbf{x}_1 - \mathbf{x}_2\|^2)$  where  $\gamma = \frac{1}{2\sigma^2}$
  - This can be expanded using Taylor series into an infinite series:

$$f(x) = f(a) + f'(a)(x-a) + \frac{f''(a)}{2!}(x-a)^2 + \frac{f^{(3)}(a)}{3!}(x-a)^3 + \dots + \frac{f^{(n)}(a)}{n!}(x-a)^n + \dots$$


$$e^x = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \frac{x^4}{4!} + \dots$$

$$\exp(-\gamma\|\mathbf{x} - \mathbf{y}\|^2) = 1 - \gamma\|\mathbf{x} - \mathbf{y}\|^2 + \frac{(\gamma\|\mathbf{x} - \mathbf{y}\|^2)^2}{2!} - \frac{(\gamma\|\mathbf{x} - \mathbf{y}\|^2)^3}{3!} + \dots$$





# RADIAL BASIS FUNCTION (RBF) SVM

- Classifier: 
$$f(x) = \sum_{i=1}^N \alpha_i y_i \exp(-\gamma \|x_1 - x_2\|^2) + b$$
- In sklearn we choose our  $\gamma$  (related to  $\sigma$ ) and  $C$  (margin softness) when building our model. There are default values.
- Data is not linearly separable in the original feature space.



# EXAMPLE PARAMS

- $\sigma = 1$ , large  $C$ ,  $\gamma = 0.5$





```
1 len(clf.support_vectors_)
2 6
3
4 clf.score(X, y)
5 1.0
```



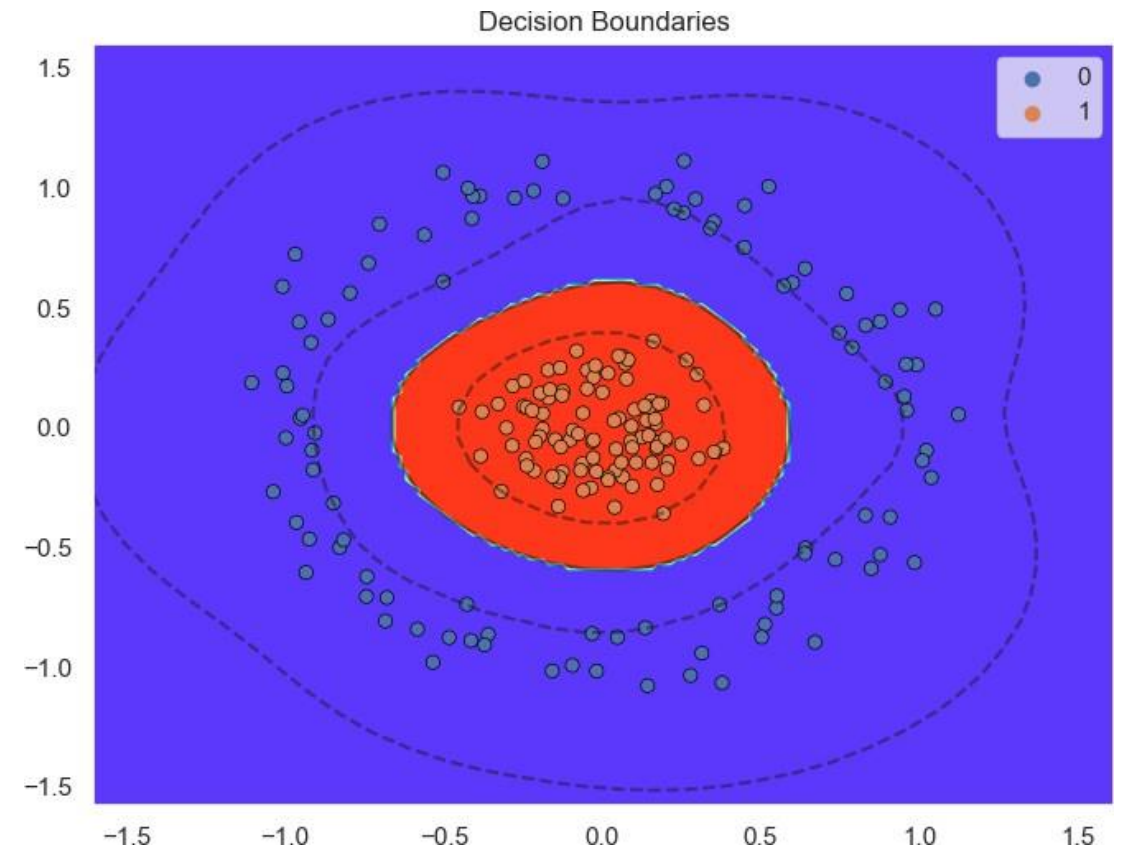

# EXAMPLE PARAMS

- $\sigma = 1, C = 0.1, \gamma = 0.5$
- Number of support vectors 171



# EXAMPLE PARAMS

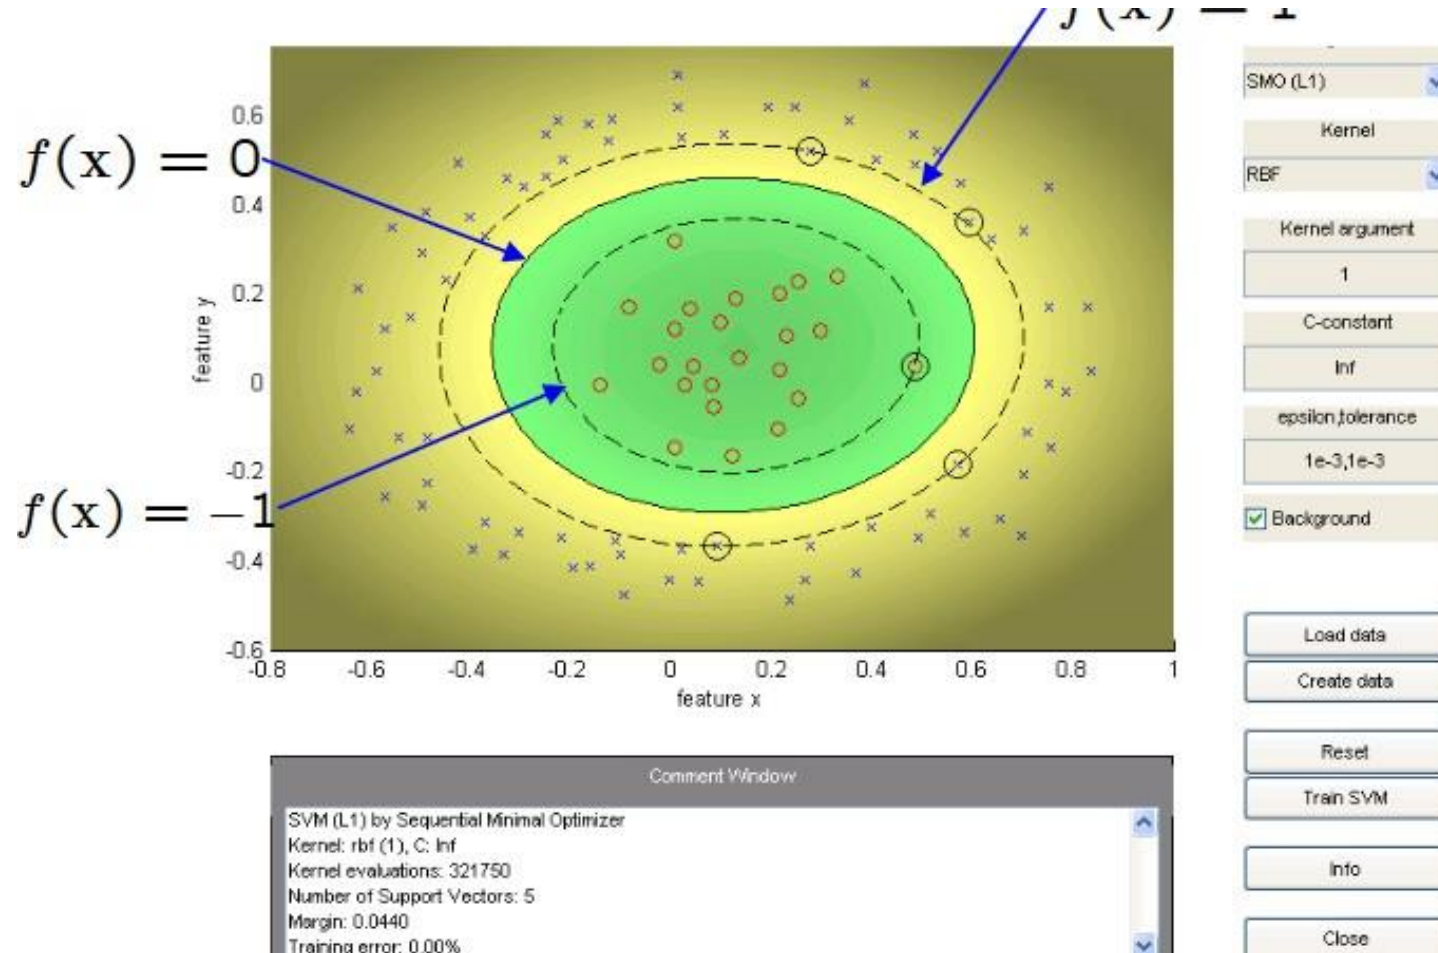
- $C = 0.1$  ,  $\gamma = 3$
- The bigger  $\gamma$  gets, the closer towards nearest neighbour classifier the rbf kernel looks.






# EXAMPLE PARAMS

- $C = \text{infinity}$ ,  $\gamma = 0.5$




# EXAMPLE PARAMS

- $C = 10, \gamma = 0.5$



# EXAMPLE PARAMS


- $C = 10, \gamma = 8$
- The bigger  $\gamma$  gets, the closer towards nearest neighbour classifier the rbf kernel looks.






# KERNEL SUMMARY


Decision boundaries of linear kernel in SVC




Decision boundaries of poly kernel in SVC



Decision boundaries of rbf kernel in SVC




Decision boundaries of sigmoid kernel in SVC



# KERNEL TRICK - SUMMARY

- Classifiers can be learnt for high dimensional features spaces, without actually having to map the points into the high dimensional space
- Data may be linearly separable in the high dimensional space, but not linearly separable in the original feature space



"I THINK YOU SHOULD BE MORE EXPLICIT  
HERE IN STEP TWO."



# SVM TIPS AND TRICKS

- SVMs are not scale invariant
- Check if your library normalizes by default
- Normalize your data
  - mean: 0 , stddev: 1
  - map to [0,1] or [-1,1]
  - StandardScaler, MinMaxScaler
- Normalize test set in same way



# PARAMETER TUNING

- Given a classification task
  - Which kernel?
  - Which kernel parameter values ( $\gamma$ )?
  - Which value for C ?
- *GridSearchCV* can be used for trying the different combinations
- By default, GridSearchCV does a 5-Fold Cross Validation test

